

# Smart Document Intelligence System - Complete Project Plan

## Project Overview

**Project Name:** Smart Document Intelligence System

**Goal:** Build a production-ready ML system that classifies documents and extracts key information

**Timeline:** 5 Sprints (5 weeks)

**Deployment:** AWS SageMaker + Web Application

**Tech Stack:** Python, PyTorch, SageMaker, Streamlit, Docker

---

## What We're Building

### Document Types (4 Classes)

1. **Invoices** - Extract: amount, vendor, date, invoice number
2. **Receipts** - Extract: merchant, total, date, items
3. **Resumes** - Extract: name, skills, experience, education
4. **Contracts** - Extract: parties, dates, key terms

### Core Features

- Upload document (PDF/Image)
  - Classify document type (98%+ accuracy)
  - Extract key information automatically
  - Display structured results
  - Export results (JSON/CSV)
- 

## Data Sources & Models

### Datasets (All FREE & Public)

1. **RVL-CDIP Dataset** - 400,000 document images, 16 classes
  - Source: [https://huggingface.co/datasets/aharley/rvl\\_cdip](https://huggingface.co/datasets/aharley/rvl_cdip)
  - Size: ~100GB (we'll use subset)
  - Classes: invoice, resume, email, memo, letter, etc.
2. **SROIE Dataset** - 1,000 receipt images with annotations
  - Source: <https://rrc.cvc.uab.es/?ch=13>

- Size: ~500MB
- Perfect for receipt extraction

### 3. Resume Dataset - Kaggle

- Source: <https://www.kaggle.com/datasets> (search "resume")
- Alternative: Generate synthetic resumes

### 4. Contract Samples - Public contracts

- Source: EDGAR SEC filings (public contracts)
- Alternative: Use RVL-CDIP contract samples

## Model Architecture

**Primary Model:** LayoutLMv3 (Microsoft)

- **Why:** State-of-the-art for document understanding (95%+ accuracy)
- **Advantages:** Understands text + layout + visual features
- **Pre-trained:** Yes (on millions of documents)
- **License:** Check HuggingFace (some have commercial restrictions)
- **Alternative:** Donut (MIT license, fully open)

**Backup/Comparison Models:**

- ResNet50 (for pure image classification)
- BERT (for text-only extraction)

## Project Structure

```
document-intelligence-ml/
├── README.md
├── requirements.txt
├── .env.example
├── .gitignore
└── setup.py
|
└── data/
    ├── raw/          # Downloaded datasets
    └── processed/   # Cleaned data
        └── samples/   # Test samples
|
└── notebooks/
```

```
|   ├── sprint1_eda.ipynb  
|   ├── sprint2_preprocessing.ipynb  
|   ├── sprint3_training.ipynb  
|   ├── sprint4_evaluation.ipynb  
|   └── sprint5_deployment.ipynb  
  
|  
|   ├── src/  
|   |   ├── __init__.py  
|   |   ├── config.py  
|   |  
|   |   ├── data/  
|   |   |   ├── __init__.py  
|   |   |   ├── download.py  
|   |   |   ├── preprocess.py  
|   |   |   └── augmentation.py  
|   |  
|   |   ├── models/  
|   |   |   ├── __init__.py  
|   |   |   ├── classifier.py  
|   |   |   ├── extractor.py  
|   |   |   └── inference.py  
|   |  
|   |   ├── training/  
|   |   |   ├── __init__.py  
|   |   |   ├── train.py  
|   |   |   ├── evaluate.py  
|   |   |   └── metrics.py  
|   |  
|   |   ├── pipeline/  
|   |   |   ├── __init__.py  
|   |   |   ├── sagemaker_pipeline.py  
|   |   |   └── step_functions.py  
|   |  
|   |   └── deployment/  
|   |       ├── __init__.py  
|   |       ├── lambda_functions/  
|   |       ├── api.py  
|   |       └── monitoring.py  
|  
|   └── web_app/  
|       ├── app.py  
|       ├── requirements.txt  
|       ├── Dockerfile  
|       └── static/  
|           ├── css/  
|           └── js/  
|           └── templates/
```

```
|- index.html  
|- results.html  
  
|- tests/  
|   |- test_data.py  
|   |- test_models.py  
|   \- test_api.py  
  
|- deployment/  
|   |- cloudformation/  
|   |- terraform/  
|   \- docker-compose.yml  
  
|- docs/  
|   \- api_documentation.md  
|   \- model_card.md  
|   \- blog_post.md
```

## 🚀 SPRINT BREAKDOWN

### SPRINT 0: SETUP & PREPARATION (Before Sprint 1)

#### Tasks

##### 1. Environment Setup (1 hour)

- Install Python 3.9+
- Install AWS CLI
- Configure AWS credentials
- Install required packages

#### Commands:

```
bash
```

```
# Create virtual environment
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate

# Install packages
pip install --upgrade pip
pip install boto3 sagemaker pandas numpy matplotlib seaborn
pip install torch torchvision transformers datasets
pip install scikit-learn opencv-python pillow pytesseract
pip install streamlit plotly jupyterlab
```

## 2. AWS Setup (1 hour)

- Create AWS account (if needed)
- Create S3 bucket: `document-intelligence-data`
- Create IAM role for SageMaker
- Set up SageMaker Studio or Notebook instance

## 3. GitHub Repository Setup (30 min)

- Create repository: `document-intelligence-ml`
- Initialize with README
- Create project structure
- Set up .gitignore

### Initial Files to Create:

#### `requirements.txt`:

```
boto3==1.28.0
sagemaker==2.180.0
torch==2.0.1
transformers==4.35.0
datasets==2.14.0
pandas==2.1.0
numpy==1.24.0
matplotlib==3.7.0
seaborn==0.12.0
scikit-learn==1.3.0
opencv-python==4.8.0
Pillow==10.0.0
streamlit==1.28.0
pytesseract==0.3.10
python-dotenv==1.0.0
```

## config.py:

```
python

import os
from pathlib import Path

# Project paths
PROJECT_ROOT = Path(__file__).parent.parent
DATA_DIR = PROJECT_ROOT / "data"
RAW_DATA_DIR = DATA_DIR / "raw"
PROCESSED_DATA_DIR = DATA_DIR / "processed"
MODEL_DIR = PROJECT_ROOT / "models"

# AWS Configuration
AWS_REGION = os.getenv("AWS_REGION", "us-east-1")
S3_BUCKET = os.getenv("S3_BUCKET", "document-intelligence-data")
SAGEMAKER_ROLE = os.getenv("SAGEMAKER_ROLE")

# Model Configuration
DOCUMENT_CLASSES = ["invoice", "receipt", "resume", "contract"]
IMAGE_SIZE = (224, 224)
BATCH_SIZE = 16
LEARNING_RATE = 2e-5
EPOCHS = 5

# Extraction Fields
EXTRACTION_CONFIG = {
    "invoice": ["invoice_number", "date", "vendor", "total_amount"],
    "receipt": ["merchant", "date", "total", "items"],
    "resume": ["name", "email", "phone", "skills", "experience"],
    "contract": ["parties", "effective_date", "expiration_date", "terms"]
}
```

# SPRINT 1: DATA COLLECTION & EXPLORATION (Week 1)

**Goal:** Understand the data, download datasets, perform EDA

**Duration:** 5 days

**Output:** Clean dataset ready for training

## Day 1: Dataset Research & Download Strategy

### Task 1.1: Research Datasets (2 hours)

- Review RVL-CDIP dataset documentation

- Review SROIE dataset
- Identify resume dataset on Kaggle
- Document dataset statistics

**Deliverable:** [docs/dataset\\_research.md](#)

## Task 1.2: Create Download Scripts (3 hours)

- Write script to download RVL-CDIP subset
- Write script to download SROIE
- Write script to download resume dataset
- Add progress tracking

**Code:** [src/data/download.py](#)

```
python

from datasets import load_dataset
import boto3
from tqdm import tqdm

def download_rvl_cdip(num_samples_per_class=1000):
    """Download subset of RVL-CDIP dataset"""
    print("Downloading RVL-CDIP dataset...")

    # Load from HuggingFace
    dataset = load_dataset("aharley/rvl_cdip", split="train")

    # Filter for our 4 classes
    target_classes = ['invoice', 'resume', 'letter', 'email']

    # Sample and save
    # ... implementation

    return dataset

def download_sroie():
    """Download SROIE receipt dataset"""
    # Implementation
    pass

def upload_to_s3(local_path, s3_path):
    """Upload data to S3"""
    s3 = boto3.client('s3')
    # Implementation
    pass
```

### **Task 1.3: Download Data (2 hours)**

- Run download scripts
- Verify data integrity
- Upload to S3
- Document dataset structure

#### **Success Criteria:**

- At least 4,000 images downloaded (1,000 per class)
  - Data uploaded to S3
  - No corrupted files
- 

## **Day 2-3: Exploratory Data Analysis**

### **Task 1.4: Create EDA Notebook (4 hours)**

- Load datasets
- Visualize sample images
- Analyze class distribution
- Check image dimensions
- Analyze file formats
- Identify data quality issues

**Notebook:** `notebooks/sprint1_eda.ipynb`

#### **Key Analyses:**

python

```

# 1. Class Distribution
import matplotlib.pyplot as plt
import seaborn as sns

class_counts = df['label'].value_counts()
plt.figure(figsize=(10, 6))
sns.barplot(x=class_counts.index, y=class_counts.values)
plt.title('Document Class Distribution')
plt.show()

# 2. Image Size Distribution
widths = []
heights = []
for img_path in tqdm(image_paths):
    img = Image.open(img_path)
    widths.append(img.width)
    heights.append(img.height)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.hist(widths, bins=50)
plt.title('Image Widths')
plt.subplot(1, 2, 2)
plt.hist(heights, bins=50)
plt.title('Image Heights')
plt.show()

# 3. Sample Visualizations
fig, axes = plt.subplots(2, 4, figsize=(20, 10))
for i, doc_class in enumerate(DOCUMENT_CLASSES):
    # Show 2 examples per class
    pass

```

## Task 1.5: Data Quality Report (2 hours)

- Document findings
- Identify preprocessing needs
- Plan data cleaning strategy

**Deliverable:** `docs/data_quality_report.md`

---

## Day 4-5: Data Preprocessing

### Task 1.6: Build Preprocessing Pipeline (6 hours)

- Create image preprocessing functions
- Handle different image formats
- Resize images
- Normalize images
- Create train/val/test splits

**Code:** `src/data/preprocess.py`

```
python
```

```

import cv2
from PIL import Image
import numpy as np

class DocumentPreprocessor:
    def __init__(self, target_size=(224, 224)):
        self.target_size = target_size

    def preprocess_image(self, image_path):
        """Preprocess single image"""
        # Load image
        img = Image.open(image_path).convert('RGB')

        # Resize
        img = img.resize(self.target_size)

        # Convert to array
        img_array = np.array(img)

        # Normalize
        img_array = img_array / 255.0

    return img_array

def create_splits(self, df, train_ratio=0.7, val_ratio=0.15):
    """Create train/val/test splits"""
    # Stratified split
    from sklearn.model_selection import train_test_split

    train_df, temp_df = train_test_split(
        df, train_size=train_ratio, stratify=df['label']
    )
    val_df, test_df = train_test_split(
        temp_df, train_size=val_ratio/(1-train_ratio),
        stratify=temp_df['label']
    )

    return train_df, val_df, test_df

```

## Task 1.7: Process All Data (4 hours)

- Run preprocessing on all images
- Save processed data
- Create data manifest files
- Upload to S3

## Task 1.8: Create Data Loaders (2 hours)

- Build PyTorch Dataset class
- Build DataLoader
- Test loading speed
- Add data augmentation

### Code Example:

```
python

from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms

class DocumentDataset(Dataset):
    def __init__(self, df, transform=None):
        self.df = df
        self.transform = transform

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        img_path = self.df.iloc[idx]['image_path']
        label = self.df.iloc[idx]['label']

        # Load image
        image = Image.open(img_path).convert('RGB')

        if self.transform:
            image = self.transform(image)

        return image, label

# Augmentation
train_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomRotation(5),
    transforms.ColorJitter(0.2, 0.2, 0.2),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])
```

## Sprint 1 Deliverables Checklist

- Dataset downloaded (4,000+ images)
- EDA notebook completed
- Data quality report
- Preprocessing pipeline built
- Train/val/test splits created
- Data uploaded to S3
- DataLoaders tested and working

### Success Metrics:

- 1,000 images per class minimum
  - Clean 70/15/15 split
  - No corrupted images
  - Average preprocessing time < 100ms per image
- 

## SPRINT 2: MODEL DEVELOPMENT (Week 2)

**Goal:** Build classification model, achieve 90%+ accuracy

**Duration:** 5 days

**Output:** Trained document classifier

### Day 1: Model Architecture Setup

#### Task 2.1: Research Model Options (2 hours)

- Review LayoutLMv3 documentation
- Check model licenses
- Compare alternatives (Donut, ResNet)
- Make architecture decision

#### Task 2.2: Setup Base Model (4 hours)

- Load pre-trained LayoutLMv3
- Modify for 4-class classification
- Test forward pass
- Calculate model size

**Code:** `src/models/classifier.py`

```
python
```

```
import torch
import torch.nn as nn
from transformers import LayoutLMv3ForSequenceClassification, LayoutLMv3Processor

class DocumentClassifier(nn.Module):
    def __init__(self, num_classes=4, model_name="microsoft/layoutlmv3-base"):
        super().__init__()

        # Load pre-trained model
        self.model = LayoutLMv3ForSequenceClassification.from_pretrained(
            model_name,
            num_labels=num_classes
        )

        self.processor = LayoutLMv3Processor.from_pretrained(model_name)

    def forward(self, images, words=None, boxes=None):
        # Process inputs
        encoding = self.processor(
            images=images,
            text=words,
            boxes=boxes,
            return_tensors="pt",
            padding="max_length",
            truncation=True
        )

        # Forward pass
        outputs = self.model(**encoding)
        return outputs

    # Alternative: Simple CNN if LayoutLM too heavy
class SimpleCNN(nn.Module):
    def __init__(self, num_classes=4):
        super().__init__()

        # Use ResNet backbone
        import torchvision.models as models
        self.backbone = models.resnet50(pretrained=True)

        # Replace final layer
        num_features = self.backbone.fc.in_features
        self.backbone.fc = nn.Linear(num_features, num_classes)

    def forward(self, x):
        return self.backbone(x)
```

---

## Day 2-3: Training Pipeline

### Task 2.3: Build Training Script (6 hours)

- Create training loop
- Add validation loop
- Implement early stopping
- Add checkpointing
- Add logging

**Code:** `src/training/train.py`

```
python
```

```
import torch
from torch.optim import AdamW
from tqdm import tqdm

class Trainer:
    def __init__(self, model, train_loader, val_loader, device='cuda'):
        self.model = model.to(device)
        self.train_loader = train_loader
        self.val_loader = val_loader
        self.device = device

    # Optimizer
    self.optimizer = AdamW(model.parameters(), lr=2e-5)

    # Loss
    self.criterion = nn.CrossEntropyLoss()

    # Best model tracking
    self.best_val_acc = 0

    def train_epoch(self):
        self.model.train()
        total_loss = 0
        correct = 0
        total = 0

        for images, labels in tqdm(self.train_loader):
            images = images.to(self.device)
            labels = labels.to(self.device)

            # Forward
            outputs = self.model(images)
            loss = self.criterion(outputs.logits, labels)

            # Backward
            self.optimizer.zero_grad()
            loss.backward()
            self.optimizer.step()

            # Metrics
            total_loss += loss.item()
            _, predicted = outputs.logits.max(1)
            correct += predicted.eq(labels).sum().item()
            total += labels.size(0)

        return total_loss / len(self.train_loader), correct / total
```

```

def validate(self):
    self.model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in self.val_loader:
            images = images.to(self.device)
            labels = labels.to(self.device)

            outputs = self.model(images)
            _, predicted = outputs.logits.max(1)
            correct += predicted.eq(labels).sum().item()
            total += labels.size(0)

    return correct / total

def fit(self, epochs=10):
    for epoch in range(epochs):
        print(f"Epoch {epoch+1}/{epochs}")

        # Train
        train_loss, train_acc = self.train_epoch()
        print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}")

        # Validate
        val_acc = self.validate()
        print(f"Val Acc: {val_acc:.4f}")

        # Save best model
        if val_acc > self.best_val_acc:
            self.best_val_acc = val_acc
            torch.save(self.model.state_dict(), 'best_model.pth')
            print("✓ Saved best model")

```

## Task 2.4: Train Initial Model (4 hours)

- Start training with baseline config
  - Monitor training metrics
  - Save training logs
  - Evaluate on validation set
-

## Day 4: Model Evaluation

### Task 2.5: Build Evaluation Script (3 hours)

- Calculate accuracy, precision, recall, F1
- Generate confusion matrix
- Analyze per-class performance
- Identify failure cases

**Code:** `src/training/evaluate.py`

```
python
```

```

from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

def evaluate_model(model, test_loader, device='cuda'):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for images, labels in test_loader:
            images = images.to(device)
            outputs = model(images)
            _, predicted = outputs.logits.max(1)

            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.numpy())

    # Classification Report
    print(classification_report(all_labels, all_preds,
                                target_names=DOCUMENT_CLASSES))

    # Confusion Matrix
    cm = confusion_matrix(all_labels, all_preds)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=DOCUMENT_CLASSES,
                yticklabels=DOCUMENT_CLASSES)
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.title('Confusion Matrix')
    plt.savefig('confusion_matrix.png')

    return all_preds, all_labels

```

## Task 2.6: Run Full Evaluation (2 hours)

- Test on test set
- Generate metrics
- Create visualizations
- Document results

## Day 5: Model Optimization

### Task 2.7: Hyperparameter Tuning (4 hours)

- Try different learning rates
- Experiment with batch sizes
- Test different architectures
- Use SageMaker automatic tuning (optional)

### Task 2.8: Model Selection (2 hours)

- Compare all model versions
  - Select best performing model
  - Save final model
  - Upload to S3 / SageMaker Model Registry
- 

## Sprint 2 Deliverables Checklist

- Document classifier trained
- Validation accuracy > 90%
- Test accuracy > 88%
- Confusion matrix generated
- Model saved to S3
- Training notebook completed
- Evaluation report written

### Success Metrics:

- Overall accuracy: > 90%
  - Per-class F1 score: > 0.85
  - Inference time: < 500ms per document
  - Model size: < 500MB
- 

## SPRINT 3: INFORMATION EXTRACTION (Week 3)

**Goal:** Build extraction models for each document type

**Duration:** 5 days

**Output:** Working extraction pipeline

## Day 1-2: OCR & Text Extraction

### Task 3.1: Setup OCR Engine (3 hours)

- Install Tesseract
- Test AWS Textract
- Compare OCR engines
- Choose best option

**Code:** `src/models/ocr.py`

```
python

import pytesseract
from PIL import Image
import boto3

class OCREngine:
    def __init__(self, engine='tesseract'):
        self.engine = engine
        if engine == 'textract':
            self.textract = boto3.client('textract')

    def extract_text(self, image_path):
        if self.engine == 'tesseract':
            img = Image.open(image_path)
            text = pytesseract.image_to_string(img)
            boxes = pytesseract.image_to_boxes(img)
            return text, boxes

        elif self.engine == 'textract':
            with open(image_path, 'rb') as document:
                response = self.textract.detect_document_text(
                    Document={'Bytes': document.read()})
            )
            return self._parse_textract_response(response)
```

### Task 3.2: Build Entity Extraction (5 hours)

- Setup spaCy or BERT NER
- Define entities for each document type
- Create extraction rules
- Test extraction accuracy

**Code:** `src/models/extractor.py`

python

```
import spacy
import re
from datetime import datetime

class InformationExtractor:
    def __init__(self):
        # Load NER model
        self.nlp = spacy.load("en_core_web_sm")

    def extract_invoice(self, text):
        """Extract invoice information"""
        results = {}

        # Invoice number (pattern: INV-XXXX)
        inv_pattern = r'INV-\d{4}'
        inv_match = re.search(inv_pattern, text, re.IGNORECASE)
        if inv_match:
            results['invoice_number'] = inv_match.group()

        # Total amount
        amount_pattern = r'\$\s*\d,?\d*'
        amounts = re.findall(amount_pattern, text)
        if amounts:
            results['total_amount'] = amounts[-1] # Last amount usually total

        # Date
        date_pattern = r'\d{1,2}[-]\d{1,2}[-]\d{2,4}'
        date_match = re.search(date_pattern, text)
        if date_match:
            results['date'] = date_match.group()

        # Vendor (using NER)
        doc = self.nlp(text)
        for ent in doc.ents:
            if ent.label_ == 'ORG':
                results['vendor'] = ent.text
                break

        return results

    def extract_receipt(self, text):
        """Extract receipt information"""
        # Similar logic for receipts
        pass

    def extract_resume(self, text):
```

```

"""Extract resume information"""
results = {}

# Email
email_pattern = r'[\w\.-]+@[ \w\.-]+\.\w+'
email_match = re.search(email_pattern, text)
if email_match:
    results['email'] = email_match.group()

# Phone
phone_pattern = r'\+?1\s*(\d{3})?(\s.-)\d{3}(\s.-)\d{4}'
phone_match = re.search(phone_pattern, text)
if phone_match:
    results['phone'] = phone_match.group()

# Skills (simple keyword matching)
skill_keywords = ['python', 'java', 'javascript', 'sql', 'aws',
                  'machine learning', 'data science']
found_skills = []
text_lower = text.lower()
for skill in skill_keywords:
    if skill in text_lower:
        found_skills.append(skill)
results['skills'] = found_skills

return results

def extract_contract(self, text):
    """Extract contract information"""
    # Similar logic for contracts
    pass

```

## Day 3-4: Extraction Pipeline Integration

### Task 3.3: Build End-to-End Pipeline (6 hours)

- Combine classifier + OCR + extractor
- Handle different document types
- Add error handling
- Test pipeline

**Code:** `src/pipeline/document_pipeline.py`

python

```

class DocumentPipeline:
    def __init__(self):
        self.classifier = DocumentClassifier.load('best_model.pth')
        self.ocr = OCREngine(engine='tesseract')
        self.extractor = InformationExtractor()

    def process_document(self, image_path):
        """Process single document end-to-end"""
        # Step 1: Classify
        doc_type, confidence = self.classifier.predict(image_path)

        # Step 2: OCR
        text, boxes = self.ocr.extract_text(image_path)

        # Step 3: Extract based on type
        if doc_type == 'invoice':
            extracted_info = self.extractor.extract_invoice(text)
        elif doc_type == 'receipt':
            extracted_info = self.extractor.extract_receipt(text)
        elif doc_type == 'resume':
            extracted_info = self.extractor.extract_resume(text)
        elif doc_type == 'contract':
            extracted_info = self.extractor.extract_contract(text)

        # Return results
        return {
            'document_type': doc_type,
            'confidence': confidence,
            'extracted_info': extracted_info,
            'full_text': text
        }

```

### Task 3.4: Test on Real Documents (4 hours)

- Test on 100 documents
  - Calculate extraction accuracy
  - Identify common errors
  - Improve extraction rules
- 

## Day 5: Validation & Optimization

### Task 3.5: Validate Extraction Quality (4 hours)

- Manual review of extractions

- Calculate precision/recall per field
- Fix critical bugs
- Document limitations

### Task 3.6: Optimize Performance (2 hours)

- Profile code
  - Optimize slow functions
  - Add caching where appropriate
- 

## Sprint 3 Deliverables Checklist

- OCR working on all document types
- Extraction models for all 4 types
- End-to-end pipeline functional
- Extraction accuracy > 80% per field
- Pipeline tested on 100+ documents
- Performance optimizations complete

### Success Metrics:

- OCR accuracy: > 95%
  - Field extraction accuracy: > 80%
  - Pipeline processing time: < 2 seconds per document
  - Error handling robust
- 

## SPRINT 4: SAGEMAKER DEPLOYMENT (Week 4)

**Goal:** Deploy models to AWS SageMaker

**Duration:** 5 days

**Output:** Production-ready ML endpoints

### Day 1: SageMaker Pipeline Setup

#### Task 4.1: Create Training Job (4 hours)

- Convert training script to SageMaker format
- Configure instance types
- Set hyperparameters
- Test training job

**Code:** `src/pipeline/sagemaker_pipeline.py`

```

python

import sagemaker
from sagemaker.pytorch import PyTorch

def create_training_job():
    sagemaker_session = sagemaker.Session()
    role = sagemaker.get_execution_role()

    # Create PyTorch estimator
    estimator = PyTorch(
        entry_point='train.py',
        source_dir='src/training',
        role=role,
        framework_version='2.0.0',
        py_version='py310',
        instance_count=1,
        instance_type='ml.p3.2xlarge', # GPU instance
        hyperparameters={

            'epochs': 10,
            'batch-size': 16,
            'learning-rate': 2e-5
        }
    )

    # Start training
    estimator.fit({'training': 's3://bucket/data/train'})

    return estimator

```

## Task 4.2: Setup Model Registry (2 hours)

- Register trained models
  - Add model metadata
  - Version models
  - Test model retrieval
- 

## Day 2: Model Deployment

### Task 4.3: Create Inference Script (3 hours)

- Write inference.py
- Handle preprocessing
- Handle postprocessing

- Test locally

**Code:** [src/deployment/inference.py](#)

```
python

import torch
import json

def model_fn(model_dir):
    """Load model"""
    model = DocumentClassifier()
    model.load_state_dict(torch.load(f'{model_dir}/model.pth'))
    return model

def input_fn(request_body, content_type):
    """Preprocess input"""
    if content_type == 'application/json':
        data = json.loads(request_body)
        # Process image
        return data
    else:
        raise ValueError(f"Unsupported content type: {content_type}")

def predict_fn(input_data, model):
    """Make prediction"""
    model.eval()
    with torch.no_grad():
        output = model(input_data)
    return output

def output_fn(prediction, accept):
    """Format output"""
    if accept == 'application/json':
        return json.dumps(prediction), accept
    raise ValueError(f"Unsupported accept type: {accept}")
```

## Task 4.4: Deploy Model Endpoint (3 hours)

- Create endpoint configuration
- Deploy endpoint
- Test endpoint
- Monitor deployment

**Code:**

```
python
```

```
from sagemaker.pytorch import PyTorchModel

def deploy_model(estimator):
    pytorch_model = PyTorchModel(
        model_data=estimator.model_data,
        role=role,
        entry_point='inference.py',
        framework_version='2.0.0',
        py_version='py310'
    )

    predictor = pytorch_model.deploy(
        instance_type='ml.m5.large',
        initial_instance_count=1,
        endpoint_name='doc-classifier-endpoint'
    )

    return predictor
```

## Day 3-4: Lambda Functions & API

### Task 4.5: Create Lambda Functions (6 hours)

- Upload handler
- Classification handler
- Extraction handler
- Result formatter

**Code:** [src/deployment/lambda\\_functions/handler.py](#)

```
python
```

```

import boto3
import json

s3 = boto3.client('s3')
sagemaker_runtime = boto3.client('sagemaker-runtime')

def lambda_handler(event, context):
    # Get uploaded file
    bucket = event['Records'][0]['s3']['bucket']['name']
    key = event['Records'][0]['s3']['object']['key']

    # Download file
    obj = s3.get_object(Bucket=bucket, Key=key)
    image_bytes = obj['Body'].read()

    # Call SageMaker endpoint
    response = sagemaker_runtime.invoke_endpoint(
        EndpointName='doc-classifier-endpoint',
        ContentType='application/x-image',
        Body=image_bytes
    )

    # Parse results
    result = json.loads(response['Body'].read())

    # Store results in DynamoDB
    dynamodb = boto3.resource('dynamodb')
    table = dynamodb.Table('DocumentResults')
    table.put_item(Item={
        'document_id': key,
        'result': result
    })

return {
    'statusCode': 200,
    'body': json.dumps(result)
}

```

## Task 4.6: Setup API Gateway (3 hours)

- Create REST API
- Configure endpoints
- Add authentication
- Test API

## **Day 5: Monitoring & Testing**

### **Task 4.7: Setup Model Monitoring (3 hours)**

- Enable SageMaker Model Monitor
- Set up CloudWatch dashboards
- Configure alerts
- Test monitoring

### **Task 4.8: Load Testing (3 hours)**

- Run load tests
  - Measure latency
  - Test auto-scaling
  - Document performance
- 

## **Sprint 4 Deliverables Checklist**

- Models deployed to SageMaker
- Endpoints responding correctly
- Lambda functions working
- API Gateway configured
- Monitoring active
- Load testing complete

### **Success Metrics:**

- Endpoint latency: < 1 second
  - API success rate: > 99%
  - Auto-scaling working
  - Monitoring dashboards live
- 

## **SPRINT 5: WEB APPLICATION & DOCUMENTATION (Week 5)**

**Goal:** Build user-facing app and complete documentation

**Duration:** 5 days

**Output:** Deployed web app + blog post

## Day 1-2: Web Application

### Task 5.1: Build Streamlit App (8 hours)

- Create UI layout
- Add file upload
- Connect to API
- Display results
- Add export functionality

**Code:** `web_app/app.py`

```
python
```

```
import streamlit as st
import requests
from PIL import Image
import json

st.set_page_config(page_title="Document Intelligence", layout="wide")

st.title("📄 Smart Document Intelligence System")
st.markdown("Upload a document to automatically classify and extract key information")

# Sidebar
st.sidebar.header("About")
st.sidebar.info("""
This system can process:
- Invoices
- Receipts
- Resumes
- Contracts
""")"

# File upload
uploaded_file = st.file_uploader(
    "Upload Document (PDF, PNG, JPG)",
    type=['pdf', 'png', 'jpg', 'jpeg']
)

if uploaded_file:
    # Display image
    col1, col2 = st.columns(2)

    with col1:
        st.subheader("Uploaded Document")
        image = Image.open(uploaded_file)
        st.image(image, use_column_width=True)

# Process button
if st.button("Analyze Document", type="primary"):
    with st.spinner("Processing..."):
        # Call API
        files = {'file': uploaded_file.getvalue()}
        response = requests.post(
            'https://api.yourdomain.com/analyze',
            files=files
        )

        results = response.json()
```

```

# Display results
with col2:
    st.subheader("Analysis Results")

    # Document type
    doc_type = results['document_type']
    confidence = results['confidence']

    st.metric("Document Type", doc_type.capitalize())
    st.metric("Confidence", f"{confidence*100:.1f}%")

    # Extracted information
    st.subheader("Extracted Information")
    extracted = results['extracted_info']

    for key, value in extracted.items():
        st.text_input(
            key.replace('_', ' ').title(),
            value,
            key=key
        )

    # Export options
    st.download_button(
        "⬇️ Download JSON",
        data=json.dumps(results, indent=2),
        file_name='document_analysis.json',
        mime='application/json'
    )

```

## Task 5.2: Add Advanced Features (4 hours)

- Batch processing
  - History view
  - Statistics dashboard
  - Settings page
- 

## Day 3: Deployment

### Task 5.3: Containerize Application (3 hours)

- Create Dockerfile
- Build container
- Test locally

- Push to ECR

#### Dockerfile:

```
dockerfile

FROM python:3.9-slim

WORKDIR /app

COPY requirements.txt .
RUN pip install -r requirements.txt

COPY ..

EXPOSE 8501

CMD ["streamlit", "run", "app.py"]
```

#### Task 5.4: Deploy to Production (3 hours)

- Deploy to AWS (ECS/EC2/Streamlit Cloud)
  - Configure domain
  - Setup HTTPS
  - Test production deployment
- 

#### Day 4-5: Documentation & Blog Post

##### Task 5.5: Write Technical Blog Post (6 hours)

- Introduction & problem statement
- Data analysis section
- Model development
- Results & evaluation
- Deployment architecture
- Lessons learned
- Future improvements

Structure: `docs/blog_post.md`

##### Task 5.6: Complete Documentation (4 hours)

- README with setup instructions
- API documentation
- Model card

- User guide
- Contributing guide

### Task 5.7: Create Demo Video (2 hours)

- Record screen demo
  - Add narration
  - Upload to YouTube
  - Add to README
- 

### Sprint 5 Deliverables Checklist

- Web application deployed
  - Application accessible via URL
  - All features working
  - Blog post published
  - Documentation complete
  - Demo video created
  - GitHub repo polished
- 

## PROJECT COMPLETION CHECKLIST

### GitHub Repository

- README.md with clear instructions
- requirements.txt
- All code properly organized
- .gitignore configured
- LICENSE file
- CONTRIBUTING.md

### Technical Deliverables

- Document classifier (90%+ accuracy)
- Information extraction models
- SageMaker deployment
- API endpoints
- Web application
- Monitoring & logging

## Documentation

- Blog post (technical audience)
- Model card
- API documentation
- Architecture diagrams
- Demo video

## Performance Metrics

- Classification accuracy > 90%
  - Extraction accuracy > 80%
  - API latency < 1 second
  - Uptime > 99%
- 

## \$ ESTIMATED COSTS

### AWS Costs (Development)

- **SageMaker Training:** \$3-5 per hour (GPU)
- **SageMaker Endpoints:** \$0.10-0.30 per hour
- **S3 Storage:** \$0.023 per GB
- **Lambda:** First 1M requests free
- **Total Development:** ~\$50-100

### AWS Costs (Production - Low Traffic)

- **SageMaker Endpoint:** ~\$50/month
- **S3:** ~\$5/month
- **Lambda:** ~\$1/month
- **CloudWatch:** ~\$5/month
- **Total Production:** ~\$60-70/month

## Cost Optimization Tips

- Use SageMaker Serverless Inference (pay per request)
- Stop endpoints when not demoing
- Use S3 lifecycle policies

- Leverage free tier where possible
- 

## TOOLS & TECHNOLOGIES

### Core Stack

- **Language:** Python 3.9+
- **ML Framework:** PyTorch 2.0
- **Transformers:** HuggingFace Transformers
- **Cloud:** AWS (SageMaker, Lambda, S3)
- **Web:** Streamlit
- **OCR:** Tesseract / AWS Textract

### Development Tools

- **IDE:** VS Code / PyCharm
- **Notebooks:** Jupyter Lab
- **Version Control:** Git / GitHub
- **Container:** Docker
- **CI/CD:** GitHub Actions (optional)

### Libraries

```
torch==2.0.1
transformers==4.35.0
datasets==2.14.0
sagemaker==2.180.0
boto3==1.28.0
streamlit==1.28.0
pytesseract==0.3.10
opencv-python==4.8.0
pandas==2.1.0
numpy==1.24.0
scikit-learn==1.3.0
matplotlib==3.7.0
seaborn==0.12.0
```

## MVP (Minimum Viable Product)

- Classifies 4 document types with 85%+ accuracy
- Extracts key fields with 75%+ accuracy
- Working web interface
- Deployed to cloud
- Basic documentation

## Target Product

- Classifies 4 document types with 90%+ accuracy
- Extracts key fields with 80%+ accuracy
- Polished web interface
- Production-ready deployment
- Comprehensive documentation
- Technical blog post

## Stretch Goals

-  Support 8+ document types
  -  Real-time processing
  -  Multi-language support
  -  Mobile app
  -  API for third-party integrations
- 

## LEARNING OUTCOMES

By completing this project, you will learn:

1. End-to-end ML project lifecycle
2. Document AI & computer vision techniques
3. AWS SageMaker deployment
4. Production ML system design
5. API development & integration
6. Technical writing & communication

---

# SUPPORT & RESOURCES

## Documentation

- HuggingFace: <https://huggingface.co/docs>
- AWS SageMaker: <https://docs.aws.amazon.com/sagemaker>
- PyTorch: <https://pytorch.org/docs>

## Communities

- HuggingFace Forums
- AWS ML Community
- Stack Overflow

## Sample Projects

- Udacity ML Examples
  - HuggingFace Spaces
  - AWS Samples GitHub
- 

# TROUBLESHOOTING

## Common Issues

1. **Out of memory:** Reduce batch size
  2. **Slow training:** Use smaller model or GPU
  3. **Low accuracy:** More data augmentation
  4. **API timeout:** Optimize inference code
  5. **High costs:** Use Serverless Inference
- 

Ready to start? Let's begin with Sprint 0!